

Twitter Sentiment Analysis: Report

Ashish Verma | IITM Module 31 Graded Mini-Project

1. Project Overview

- **Goal:** Classify 73K tweets into Positive/Negative/Neutral using BiLSTM RNN
- **Dataset:** Twitter training CSV (latin-1 encoded) → **69,491 cleaned tweets**
- **3-class problem:** Pos/Neg/Neu (Irrelevant → Neutral mapping)

Raw → Cleaned: 73,996 → 69,491 tweets (6% duplicates removed)

2. Key EDA Findings

- **Sentiment Distribution:**

- Negative: 21,166 (30%) ← MOST COMMON
- Positive: 19,067 (27%)
- Neutral: 17,042 (25%)

Insight: Negative bias but balanced enough for training.

- **Tweet Patterns:**

Class	Top Words	Avg Length
Positive	love, great, awesome	55 chars
Negative	worst, bad, hate	60 chars (complaints)
Neutral	game, like, get	50 chars

- **Visuals:** Bar plots + word clouds show clear linguistic separation

3. RNN Methodology

- **Pipeline:**

1. Clean: Remove URLs/@mentions/#hashtags → Stemming → Stopwords
2. Tokenize: Vocab=10K → Pad to 50 tokens
3. Model: Embedding (128) → BiLSTM(64) → Dense → Softmax

- **Architecture** (1.42M params):

Embedding(10K→128)

↓

BiLSTM(64) + Dropout(0.5)

↓

BiLSTM(64) + Dropout(0.5)

↓

Dense(64, relu) → Dense(3, softmax)

- **Training:** 80/20 split | 5 epochs | Adam | Batch=64

4. Results: 88.1% TEST ACCURACY

F1-Scores: Pos(0.86) | Neg(0.89) | Neu(0.89)

Macro Avg: 0.88 ← EXCELLENT!

Test Set: 13,899 samples

Train: 55,592 samples

- **Training Curves:** Clean convergence, no overfitting

5. Challenges & Solutions

Challenge	Solution
NLTK stopwords	Cached after 1st download

Challenge	Solution
Variable lengths	<code>pad_sequences(maxlen=50)</code>
Class imbalance	Categorical encoding + LSTM
Memory (69K×50)	Vocab limit=10K

- **Key Fix:** BiLSTM (vs uni-directional) → +3% accuracy boost

6. Live Demo: Sentiment Classifier

"love borderlands great game!!"

→  POSITIVE (89% conf)

"worst game ever total waste"

→  NEGATIVE (92% conf)

"getting on borderlands with friends"

→  NEUTRAL (87% conf)

"amazon delivery late again!!"

→  NEGATIVE (90% conf)

7. Business Value

- ✓ Brand monitoring (detect frustration early)
- ✓ Customer service prioritization
- ✓ Product launch sentiment tracking
- ✓ Competitor analysis

- **Production Ready:** 1.4M params → Fast inference (~50ms/tweet)

8. Key Takeaways

- **88% accuracy** beats random baseline (33%)
- Negative tweets use "worst/bad/hate" (actionable insight)
- BiLSTM captures tweet context perfectly
- **Complete pipeline:** EDA → Model → Deploy-ready